



*Smart Platform for Buying, Selling, and Renting
Cars and Real Estate (Mobile & Website)*




Presented by:


Muhammad Zaid Al-Nahas

Table of Contents

Data Preparation for Modeling



Content-Based Recommendation System (Embedding Approach)





6

Introduction

**Today, we'll discuss building
an AI-powered
recommendation system for
TrustAsset, a real estate and
vehicle marketplace.**





7

Goal

Our goal is to enhance user experience and engagement by providing personalized recommendations. We're focusing on a content-based approach, which uses item features to find relevant suggestions.





8

Dataset

Contains detailed information about residential properties listed or recently sold across the United States, including price, physical characteristics (beds, baths, size, lot size), and location data.

Year	Size	Source	Domain	Dataset Name
2024	Over 2.2 million property listings	Kaggle	Real Estate	USA Real Estate Listings

Design and Train Neural Network for Embeddings



10

Autoencoder Model

- **A type of artificial neural network used for learning efficient data codings in an unsupervised manner. It learns to compress data from the input layer into a short code, and then uncompress that code into an output that is as close as possible to the original input.**
- **Structure: Consists of two main parts: an Encoder and a Decoder.**



11

Encoder: Compresses input features

- **Input Layer (Dimension: INPUT_DIM)**
- **Dense(256) -> ReLU, Batch Norm, Dropout(0.3)**
- **Dense(128) -> ReLU, Batch Norm, Dropout(0.3)**
- **Bottleneck (Embedding Layer): Learned representation**
- **Dense(10) -> ReLU activation (Initial Baseline)**
- **Dimension: EMBEDDING_DIM (10)**



12

Decoder: Reconstructs original features

- **Dense(128) -> ReLU, Batch Norm, Dropout(0.3)**
- **Dense(256) -> ReLU, Batch Norm, Dropout(0.3)**
- **Output Layer (Dimension: INPUT_DIM) -> Linear activation**



13

Training the Autoencoder

- **Goal:** Train the entire Autoencoder network.
- **Objective:** Minimize the difference between the original input data and the reconstructed output data.
- **Loss Function:** Mean Squared Error (MSE) is commonly used.



14


Key Findings from Experiments

- **Experimentation: Varied embedding dimension, activation, and layer sizes.**
- **Best Performance: Experiment 4 (Wider layers, Linear embedding activation, Dropout 0.2).**
- **Achieved the lowest reconstruction error (MSE/MAE).**
- **Linear activation in embedding layer was beneficial for standardized data.**

Autoencoder Model Experiments Comparison

Experiment	Embedding Dim	Embedding Activation	Hidden Layers (Encoder)	Dropout Rate	Best Val MSE	Best Val MAE	Best Epoch	Training Epochs
Baseline	10	'relu'	(256, 128)	0.3	0.0064	0.0431	51	66
Experiment 1	16	'relu'	(256, 128)	0.3	0.0067	0.0588	11	26
Experiment 2	10	'linear'	(256, 128)	0.3	0.0048	0.0408	67	82
Experiment 3	10	'linear'	(512, 256)	0.3	0.0016	0.0236	66	81
Experiment 4 (Best)	10	'linear'	(512, 256)	0.2	0.0009	0.0176	65	80

Generate Item Embeddings





18

Generate Item Embeddings

- **Goal: Obtain the learned numerical representation (embedding) for every property.**
- **Tool: Use the trained Encoder model.**
- **The Encoder is the first half of the Autoencoder we trained.**



19

Process: Passing Data Through the Encoder

- **Input: The item_features_matrix containing the preprocessed features for all 300k properties.**
- **This matrix is fed into the trained Encoder model.**
- **The Encoder applies the learned weights and biases to transform the input features.**



20

Output: The Item Embeddings

- **Result: The output of the Encoder for each property's feature vector is its corresponding item embedding.**
- **This is a dense, lower-dimensional vector (10).**
- **Each embedding vector is a unique "fingerprint" representing the property's content in the learned space.**


Calculate Embedding Similarity

Goal: Find which properties have similar embeddings.

Method: Compute Cosine Similarity between all pairs of embeddings.

Challenge: Full similarity matrix is too large (300k x 300k).

Solution: Chunked calculation – compute similarity in batches and store only the Top N (e.g., 50) most similar pairs per item.



Create Embedding Similarity Lookup


Input: Saved Top N embedding similarity pairs (e.g., in Parquet file).

Structure: Create an in-memory dictionary.

Keys: Source item index.

Values: List of (similar item index, similarity score) tuples.

Purpose: Enables very fast retrieval of similar items for any given property.



Generate Synthetic User Interaction Data

Need: Real user data is unavailable for development.

Solution: Simulate realistic user interactions (views, favorites).

Method: Use LLMs or rule-based logic based on user archetypes and listing features/similarity.

Output: A dataset of simulated user events (who did what, to which item, when).

Alignment: Simulate interactions that match backend schema concepts (Clicks, Favorites, Search Queries, Preferences).



27

Why Simulate Interactions?

Develop & Test: Build and refine the recommendation logic *before* live data integration.

Evaluate Personalization: Assess if the system recommends items similar to a user's *simulated* history.

Mimic Behavior: Create patterns (e.g., browsing similar items) more realistic than random data.

Controlled Environment: Test specific scenarios and user types.

Automated Property Description



Creating unique, compelling descriptions for many listings is time-consuming.
Manual process can lead to inconsistencies in style, tone, and detail.
Scalability issues with a large volume of properties.
Need for a more efficient and consistent approach.



- **Goal:** Generate descriptions automatically from structured data.
- **Model:** Started with deepseek-ai/DeepSeek-R1-Distill-Qwen-1.5B (a general-purpose LLM).
- **Process:**
 - Load structured property data (e.g., price, beds, baths, location).
 - Format data into a text "Prompt" for the LLM.
 - LLM generates a description based on its general training.
- **Outcome:** Successfully generated descriptions, but the style was generic.




Pre-trained models are generalists.

Real estate descriptions require a specific style, tone, and focus.

Fine-tuning: Adapting a pre-trained model to a specific task/domain using a smaller, relevant dataset.

Objective: Teach the LLM to generate descriptions in the desired "Trustasset style."





- **Dataset: Small (10 properties), manually crafted (Prompt, Desired Description).**
- **Tools: Hugging Face transformers (Trainer API).**
- **Process: Load model, tokenize data, configure, train.**





- **LLM fine-tuning is memory-intensive.**
- **OutOfMemoryError encountered.**
- **1.5B parameter model was too large for direct training on available GPU.**



- **Solutions: Gradient Accumulation:** Simulate larger batches with less memory.
- **Quantization (BitsAndBytes):** Load model in lower precision (4-bit) to reduce size.
- **LoRA (Low-Rank Adaptation):** Train only small adapter layers for efficiency.

Pre-trained vs. Fine-tuned Model

Comparison (Property Description Generation)

Aspect	Base Pre-trained Model (deepseek-ai/DeepSeek-R1-Distill-Qwen-1.5B)	Fine-tuned Model (with LoRA adapters)
Training Data	Massive general internet text	Massive general internet text + Small (10 examples) task-specific data
Model Parameters	All 1.5 Billion parameters	All 1.5 Billion parameters (base frozen) + Small LoRA adapter parameters (trainable)
Training Process	Extensive pre-training by model developers	Additional training on a small dataset using PEFT (LoRA)
Memory Requirement	High (for loading and inference)	Lower (for training/loading due to LoRA & quantization)
Adaptation	None (general purpose)	Adapted to the specific task/style of the fine-tuning data
Observed Output (from comparison script)	Generic, sometimes repetitive, lacks specific real estate style.	Unexpectedly poor, repetitive, sometimes nonsensical output.
Reason for Observed Output	Trained on general text, no specific real estate domain knowledge.	Severe overfitting and lack of generalization due to extremely limited fine-tuning data (only 10 examples).



**Thank
you**