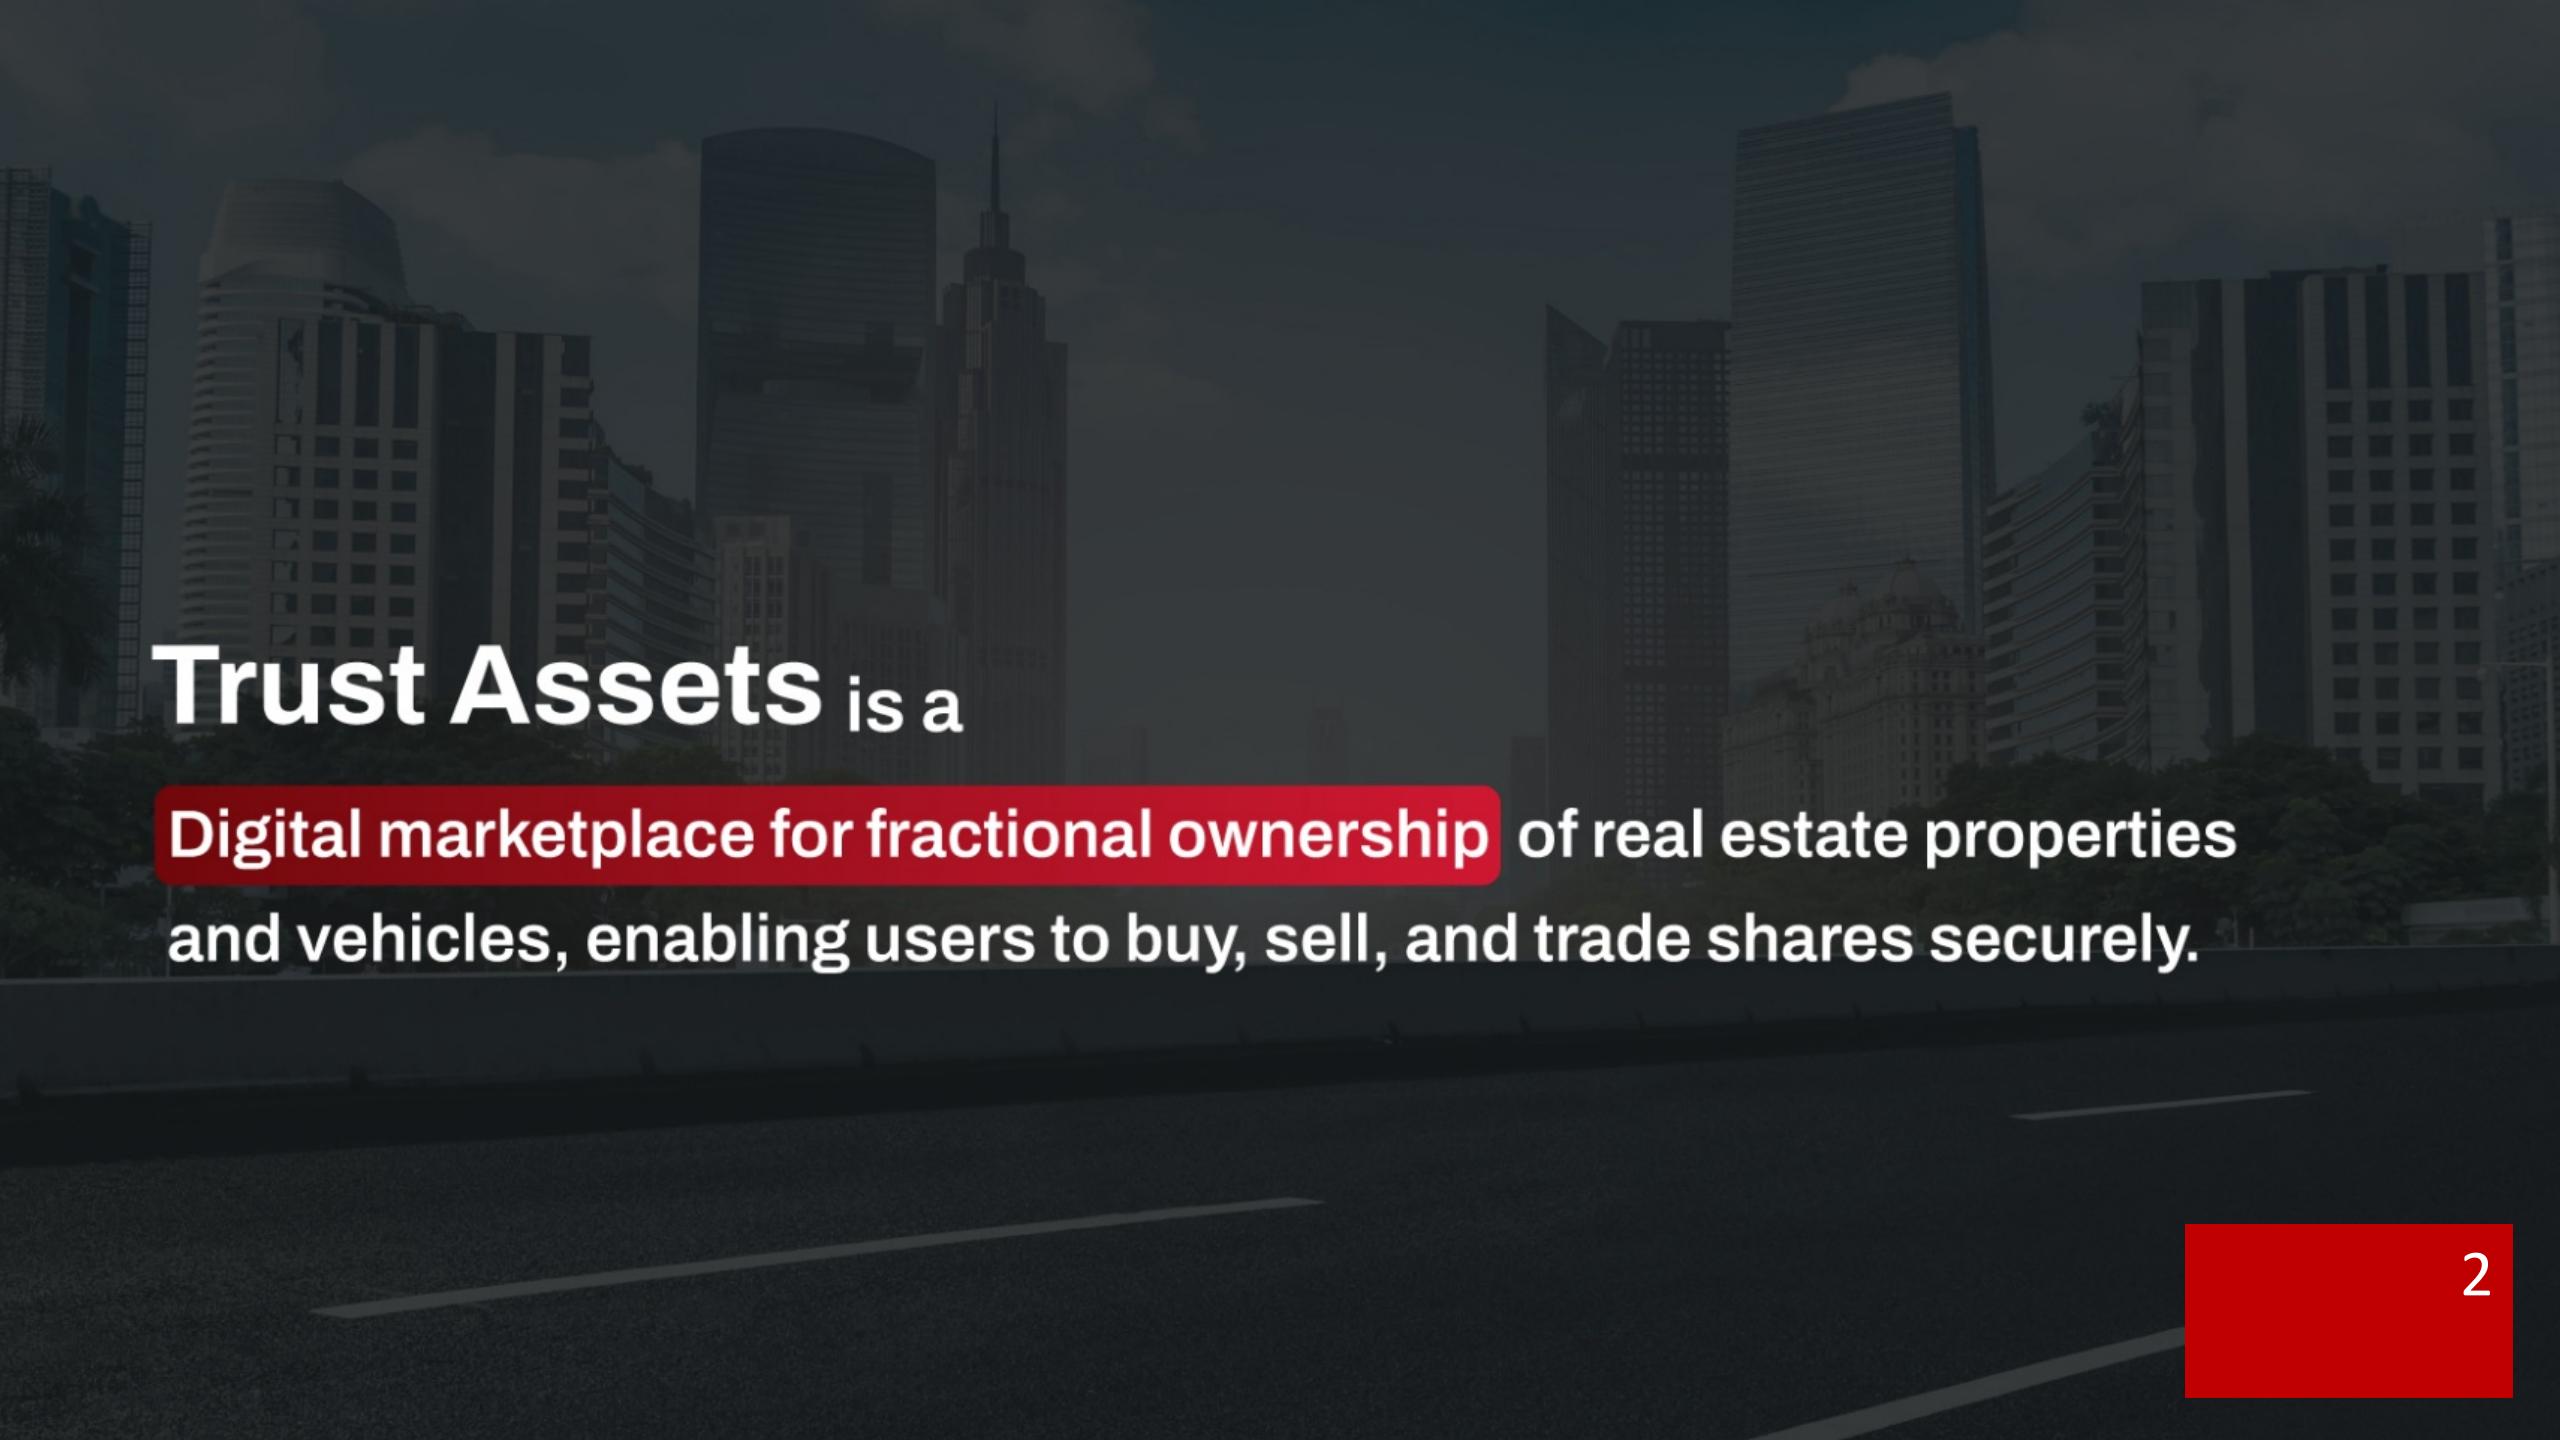




TRUST ASSETS



Trust Assets is a

Digital marketplace for fractional ownership of real estate properties and vehicles, enabling users to buy, sell, and trade shares securely.

Recommendation System

Reference	Year	Data	Preprocessing	Model	Result
(Muñoz, Meza, & Ventura, 2025)	2025	Spatial distances (health, education, security, etc.) + housing characteristics (price, size, age)	Normalization, logarithmic transformation, factor analysis, dimensionality reduction (PCA)	Fuzzy C-Means (FCM) Clustering	Identified 4 clusters (Jambu's elbow method); improved segmentation accuracy.
		Geospatial data (QGIS) + web-scraped housing features	Spatial distance calculation, variable transformation	Bat Algorithm (Metaheuristic Optimization)	Execution time: 1.38–1.42 seconds (fastest); reduced search time by 15% vs. PSO/GA.
		User preferences (sliders) + housing clusters	Cosine similarity, graph-based memory	Hybrid Recommender System (Content-Based + Collaborative Filtering)	Precision: 81.6% (improved relevance), Recall: 89% (retrieved 89% relevant items).
		Spatial factors (latent variables) + user ratings	Fuzzy membership assignment, ANOVA validation	Spatial Fuzzy Logic Model	User Satisfaction: 58% (2.9/5), highlighting need for collaborative filtering

Reference	Year	Data	Preprocessing	Model	Result
(C, Oberoi, Goyal, & Sikka, 2024)	2024	Real-world data (3M users, 1.1M properties)	Cohort creation (locality, price/area bins), scoring (leads, conversions).	Rule-Based Engine	Latency <40 ms (Ensured real-time cold-start recommendations)
			Binning (price, area), activity weighting (Table 1), binary encoding.	Content Filtering	MAP@6 = 0.881 (Buy) (Improved short-term precision by 3.2%)
			Interaction scoring (CRF=10, scrolling=1), ALS training (daily).	Collaborative Filtering (ALS)	NDCG = 0.685 (Buy) (Optimized ranking for long-term preferences)
			Combined preprocessing (content + collaborative steps).	Hybrid Model	Latency = 29.3 ms (Balanced accuracy + speed for hybrid users)

Reference	Year	Data	Preprocessing	Model	Result
(Arif, Muhtarom, & Nurhayati, 2023)	2023	50 questionnaire entries (14 criteria)		Cosine-based similarity	63.8%
				Adjusted Cosine similarity	70.4%
				Pearson correlation	88.7%
				Spearman correlation	75.57%
(Boteju & Munasinghe, 2020)	2020	Hybrid Recommender: - User data (1,873 survey records) - Item data (98 vehicle models)	- Null value removal - Categorical-to-numeric conversion - Feature engineering	Multiclass Neural Network (1 hidden layer, 1,000 nodes, 160 iterations)	96% accuracy (Improved recommendation accuracy by capturing linear/non-linear data patterns)
		NLP Sentiment Analysis: - 14,000 labeled Twitter reviews	- Tokenization - Lemmatization - Stop-word removal - Text cleaning	Naive Bayes Classifier	92.73% test accuracy (Improved sentiment classification via labeled tweet analysis)

Reference	Year	Data	Preprocessing	Model	Result
(Zhang, et al., 2019)	2019	Data1: 1-week sale data (Sydney)	Filtered items with <2 interactions	Stage 1 (Content-based)	Precision@5: 1.65% Improved baseline (1.03%) by 60%
			One-hot encoding for categorical/numerical features		
			Imputed missing prices		
		Data2: 2-week sale data (Sydney) Scenario 1/2	Same as Data1		Precision@5: 3.75% Improved baseline (2.28%) by 64.5%
			Time-based split for Scenario 2		
		Data2: 2-week sale data (Sydney) Scenario 3 (Cold-start)	Training/test split to exclude new items in training	Full two-stage model	Recall@50: 4.71% Baseline failed (0% recall)
		Data3: 3-month data (user attributes)	Accumulated user features (e.g., persona, engagement level)	Stage 2 (XGBoost spam filtering)	Precision@5: Improved from 3.75% to 4.2% Reduced spam notifications
			One-hot encoding		

Intelligent Real Estate Recommendations

A Content-Based System Powered by Deep
Learning Embeddings



1

The Challenge: Finding the Perfect Property

Our Goal

To create a smart system that understands the nuanced features of a property and recommends truly similar listings, saving users time and effort.





2

Our Two-Phase System Architecture

Offline Training

This is where we process data and train our model periodically. This is the "learning" phase.

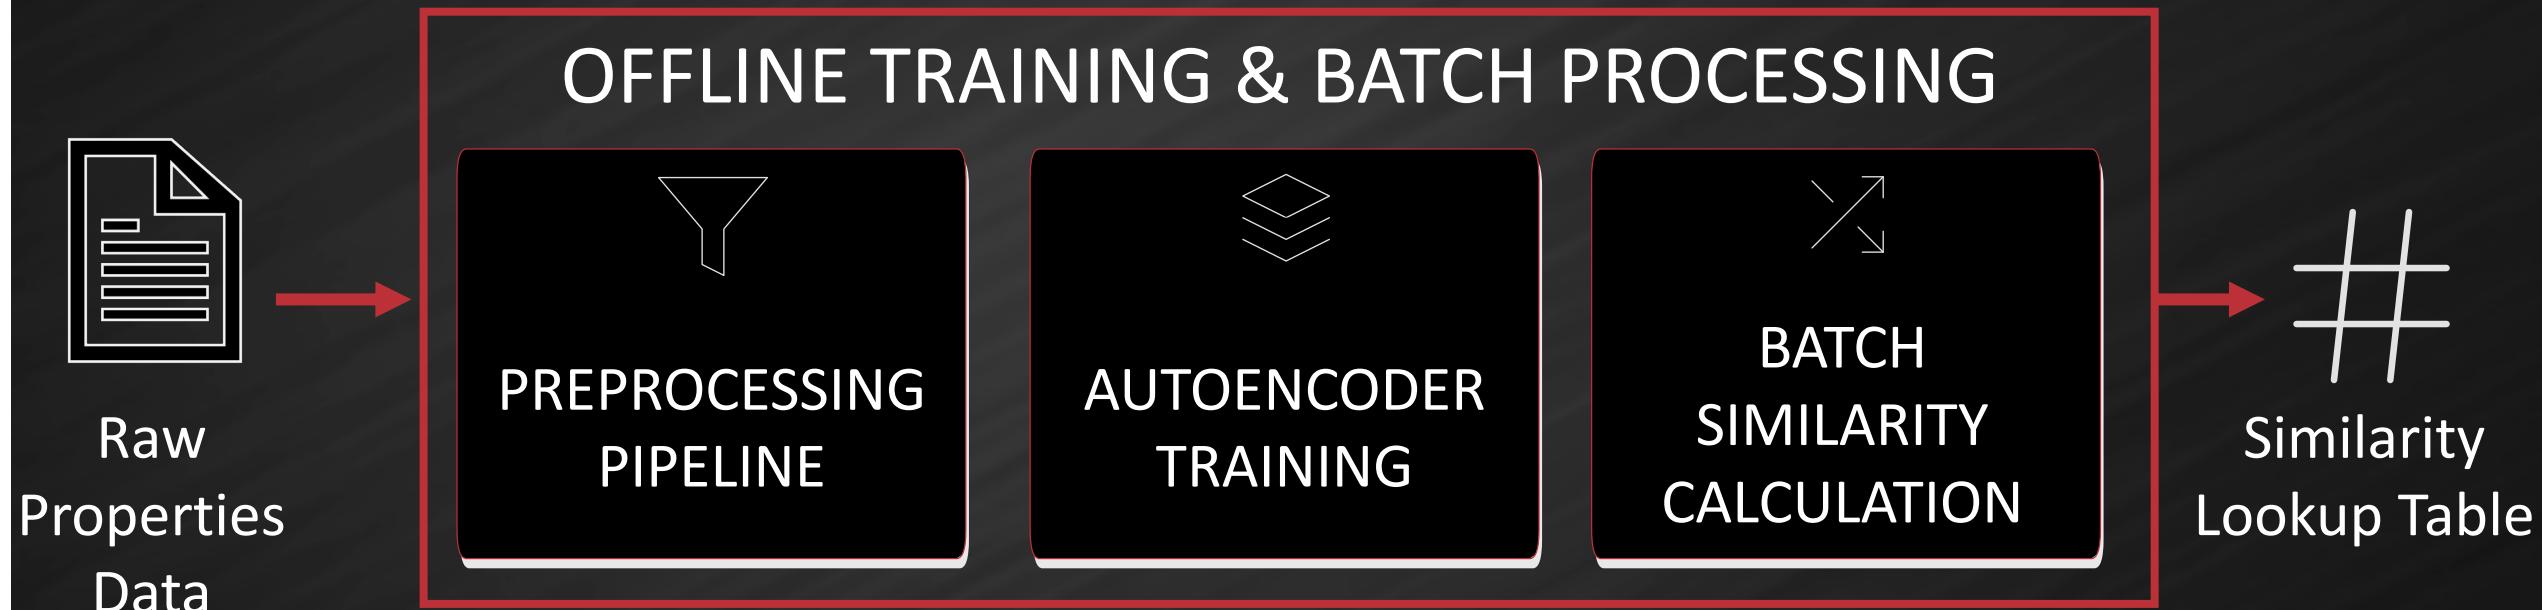
Online Serving

This is the live API that uses the pre-trained model to deliver real-time recommendations to our web and mobile apps.



3

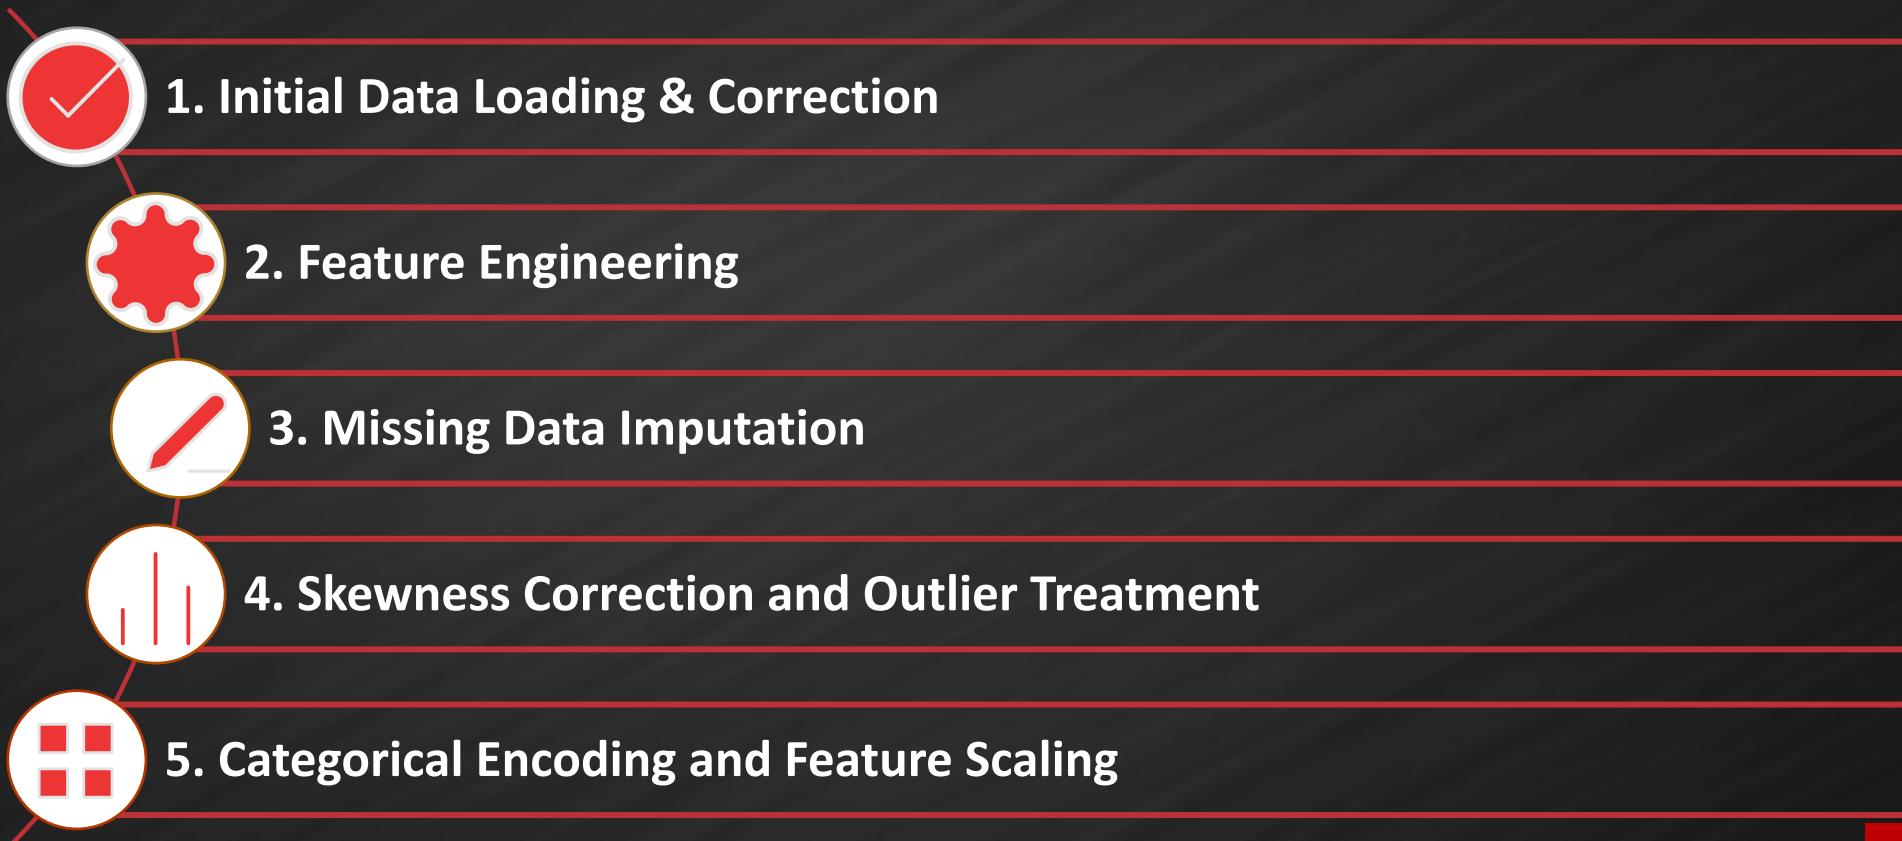
The Offline Engine - Building the Recommendation Intelligence





4

Step 1 - Preprocessing Pipeline

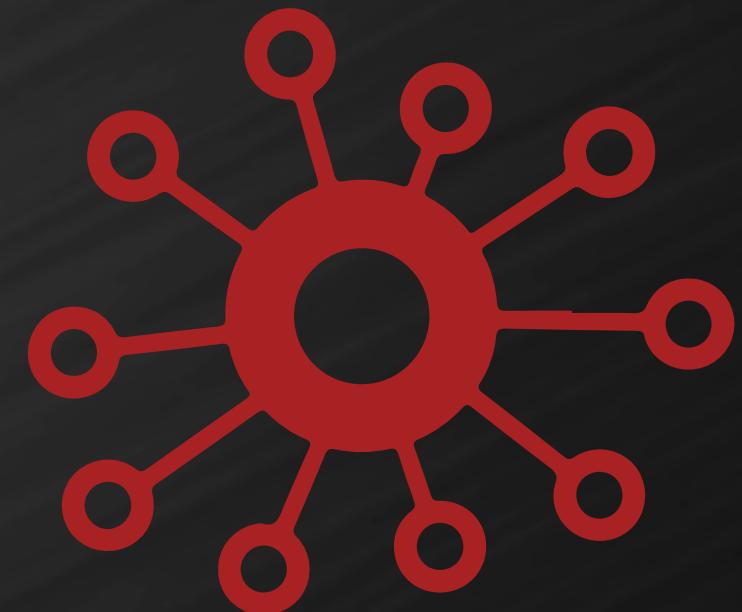




5

Step 2: Learning Property "DNA" with an Autoencoder

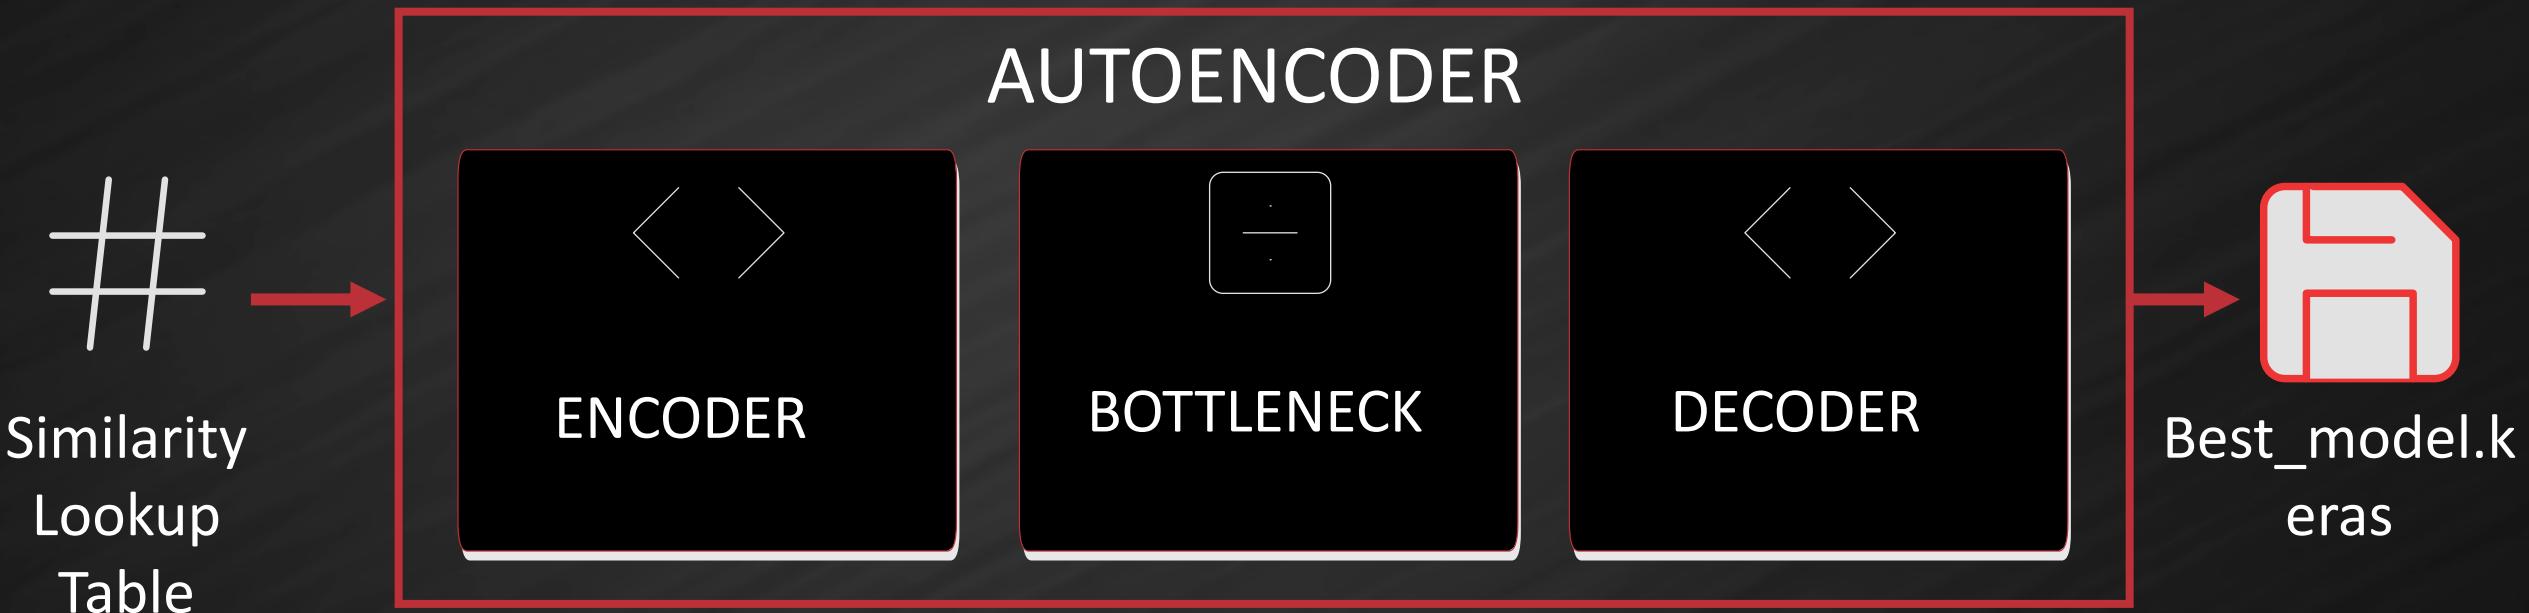
Our system uses an Autoencoder neural network to create a unique "fingerprint" for each property. This fingerprint, called an embedding, captures the property's most important features by compressing them into a dense vector.





5

Step 2: Learning Property "DNA" with an Autoencoder





5

Model Optimization & Experiments

We conducted several experiments to find the optimal model architecture.

Experiment	Embedding Dim	Embedding Activation	Hidden Layers (Neurons)	Dropout Rate	Best Val MSE
Baseline	10	'relu'	(256, 128)	0.3	0.0064
Exp 1	16	'relu'	(256, 128)	0.3	0.0067
Exp 2	10	'linear'	(256, 128)	0.3	0.0048
Exp 3	10	'linear'	(512, 256)	0.3	0.0016
Exp 4 (Final)	10	'linear'	(512, 256)	0.2	0.0009



7

Creating the Similarity Cache

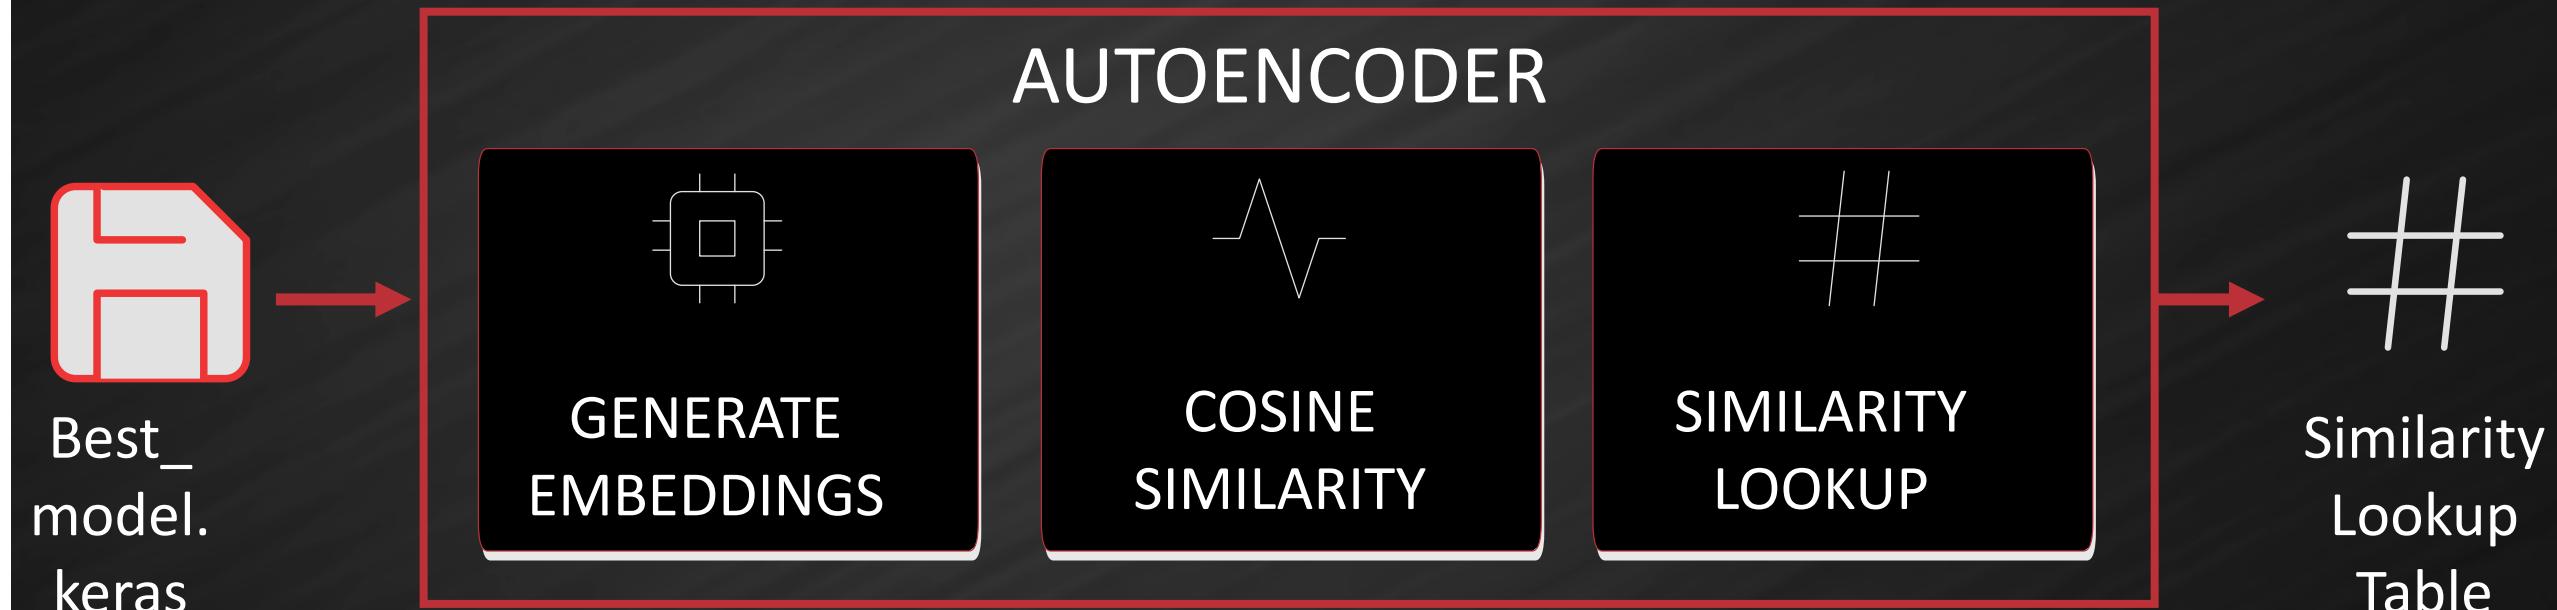


1. Generate property "fingerprints" (embeddings).
2. Calculate similarity between all fingerprints.
3. Store top 50 similar properties for quick lookup.



7

Creating the Similarity Cache





8

Phase 2: The Live Recommendation API (FastAPI)

On startup, the API loads all offline artifacts (model, pipeline, similarity lookup) into memory for maximum speed.

Exposes simple, well-defined endpoints for getting recommendations.

The screenshot shows a FastAPI documentation interface with a red border. It lists three main sections: Utilities, Recommendation, and Admin, each containing one or more API endpoints with their methods, paths, and descriptions.

Section	Method	Path	Description
Utilities	POST	/generate-embedding	Generate Embedding
Recommendation	GET	/get-similar-properties/{property_id}	Get Similar Properties
	POST	/find-similar-to-new-property	Find Similar To New Property
Admin	POST	/update-similarities	Trigger Similarity Update



8

How the API Delivers Recommendations

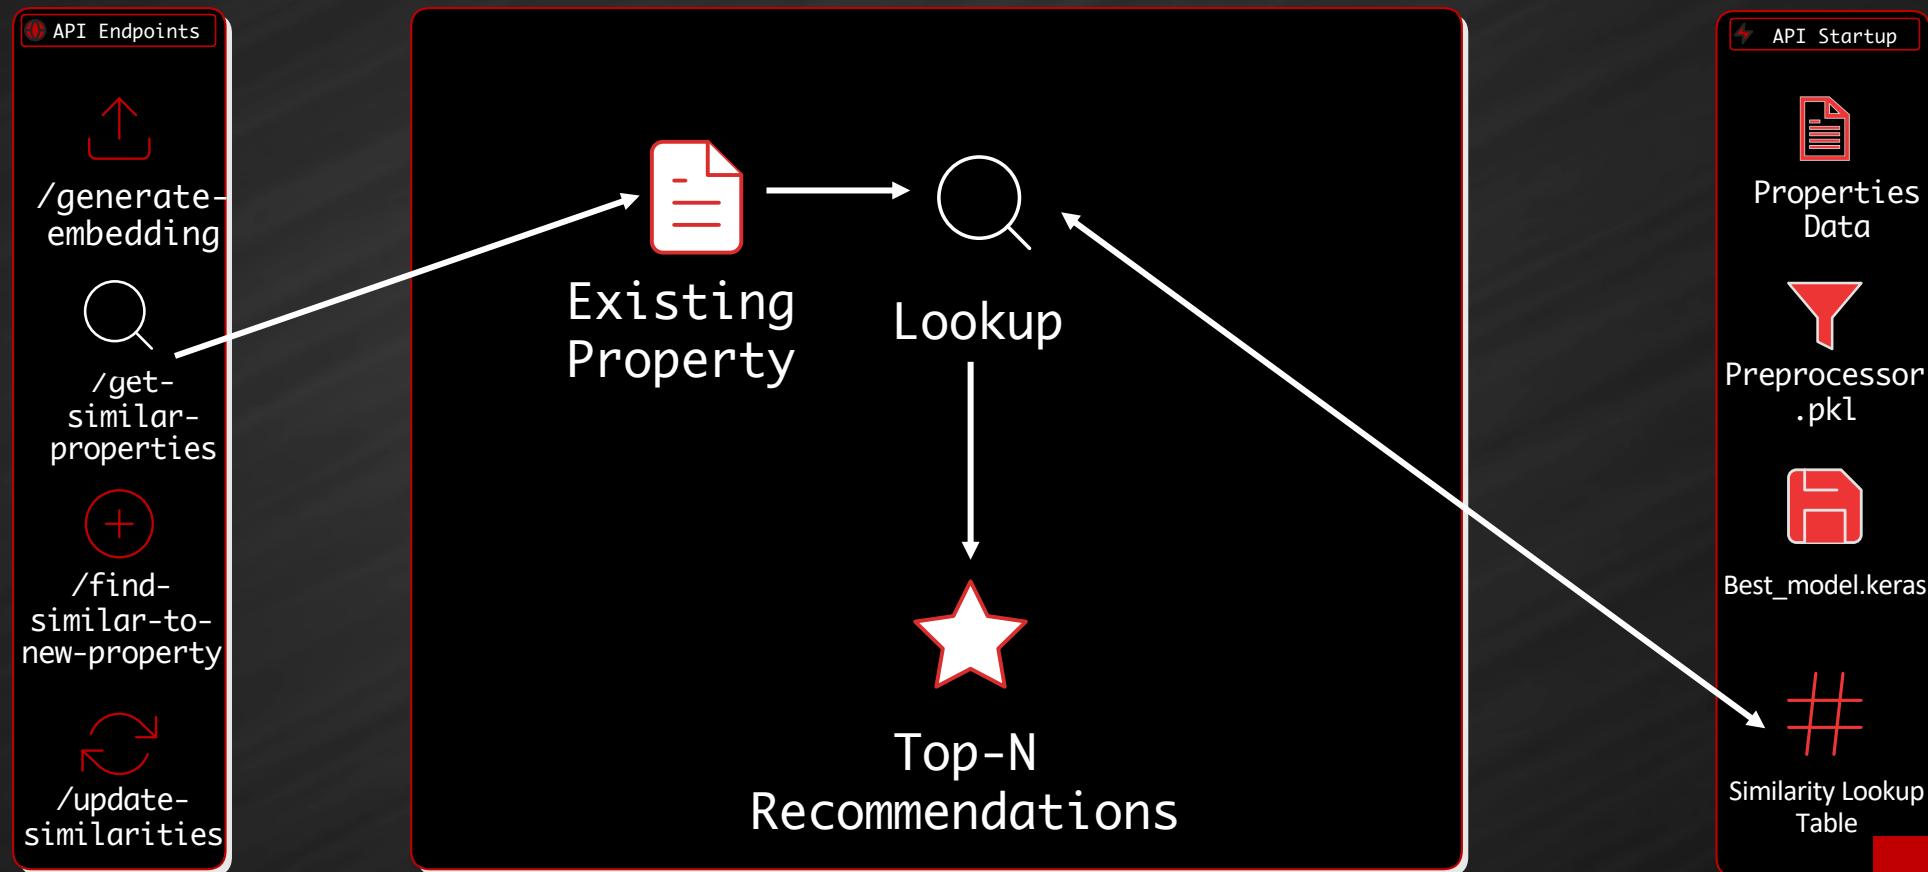
Scenario 1: Existing Property

1. Receives a property_id.
2. Performs an instant lookup in the cached **Similarity Table**.
3. Returns the pre-calculated top N similar items



8

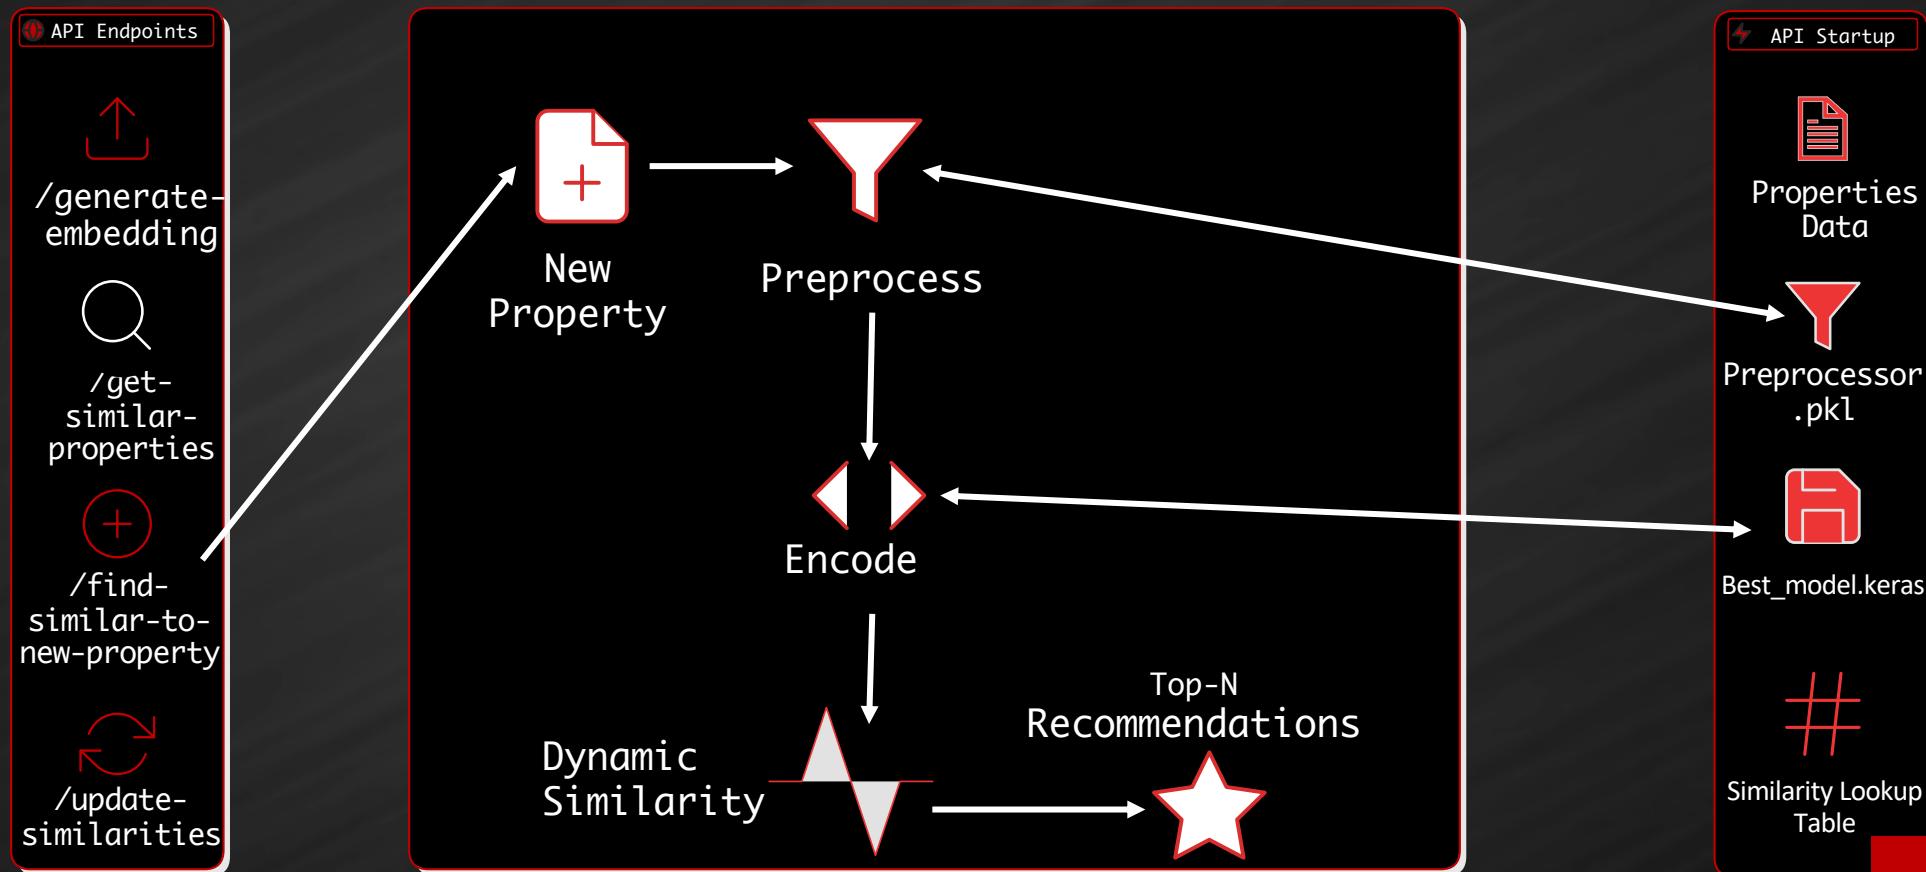
Scenario 1: Existing Property





8

Scenario 2: New Property





9

Result

We built an end-to-end system that uses deep learning embeddings to provide highly relevant, content-based property recommendations.



Vehicle Recommendation System

Powering Personalized Discovery with
Collaborative and Content-Based Filtering



1

The Goal: Increase Sales by Enhancing User Experience



Problem: Our massive inventory is overwhelming for users.

Goal: Guide users to the right car, boosting satisfaction and sales.



2

Our Solution: A Smart Recommendation Engine

Initial Methodology

- Item-Item **Collaborative Filtering**

Why

- It's a powerful, industry-standard technique that leverages community behavior (user ratings) to find similar items.

The Plan

- Analyze user ratings to understand which vehicles are perceived as similar by the community.



3

Data: The Fuel for Our Engine

Data Sources:

Domain	Name	Description	Source	Size	Year
Vehicle	Vehicle Sales Data	US used vehicle sales data including price, condition, seller, and key vehicle attributes (make, model, etc.).	Kaggle	558,837 records	2024
Vehicle	Vehicle Reviews Data	Consolidated user reviews and ratings for various vehicle models.	Kaggle	284,715 records	2018



4

Preprocessing Steps

Loaded and consolidated all data.

Handled missing values by dropping incomplete user-item-rating records.

Parsed unstructured text (Vehicle_Title) to create a standardized item_id.

Result A clean, model-ready dataset of user-item interactions



5

Model V1 & The Sparsity Challenge

Initial Model

A k-Nearest Neighbors (k-NN) model trained on a user-item rating matrix.

The Problem

Initial recommendations were irrelevant and illogical.

Root Cause: Extreme Data Sparsity (99.989%)

The model couldn't find meaningful patterns because too few users had rated the same set of cars.



6

Model V2: Addressing Sparsity by Filtering

Filter by Item Popularity
Kept only vehicles with 10+ ratings



Filter by User Activity
Kept only users with 10+ ratings.



A much smaller, but significantly denser and higher-quality dataset for model training



7

The Pivot: Evolving to a Hybrid System

The Limitation

- Pure collaborative filtering fails for new items with no ratings (the "cold start" problem).

The Solution

- Evolve the system into a **Hybrid Model**.

New Component

- A **Content-Based** recommendation engine was added to the API.

Capability

- The new engine can recommend similar vehicles based on their features (price, make, year) alone, solving the cold start problem.



8

API Endpoints Usage

GET /recommend?item_id=<ITEM_ID>

- **Method:** Collaborative Filtering
- **Description:** Retrieves pre-calculated recommendations for a known vehicle from the similarity cache.

POST /recommend/unseen

- **Method:** Content-Based Filtering
- **Description:** Generates recommendations dynamically for a new, unseen vehicle based on its features.
- **Request Body:** JSON object with vehicle features (Year, Make, Model, Trim, HorsePower, Price.).



9

Result

We successfully built and deployed a secure, memory-optimized, hybrid recommendation API that solves the cold-start problem.